

Toward less synchronous composable multilevel methods for implicit multiphysics simulation

Jed Brown¹, Mark Adams², Peter Brune¹, Matt Knepley³, Dave May⁴, Barry Smith¹

¹Mathematics and Computer Science Division, Argonne National Laboratory

²Columbia University

³Computation Institute, University of Chicago

⁴ETH Zürich

2012-02-06

Outline

Multiphysics and methods

Coupling software in PETSc

Hardware and consequences

Multiphysics problems

Examples

- ▶ Saddle-point problems (e.g. incompressibility, contact)
- ▶ Stiff waves (e.g. low-Mach combustion)
- ▶ Mixed type (e.g. radiation hydrodynamics, ALE free-surface flows)
- ▶ Multi-domain problems (e.g. fluid-structure interaction)
- ▶ Full space PDE-constrained optimization

Software/algorithmic considerations

- ▶ Separate groups develop different “physics” components
- ▶ Do not know a priori which methods will have good algorithmic properties
- ▶ Achieving high throughput is more complicated
- ▶ Multiple time and/or spatial scales
 - ▶ Splitting methods are delicate, often not in asymptotic regime
 - ▶ Strongest nonlinearities usually non-stiff: prefer explicit for TVD limiters/shocks

The Great Solver Schism: Monolithic or Split?

Split

Monolithic

- ▶ Direct solvers
- ▶ Coupled Schwarz
- ▶ Coupled Neumann-Neumann
(need unassembled matrices)
- ▶ Coupled multigrid
- ✗ Need to understand local spectral and compatibility properties of the coupled system

- ▶ Physics-split Schwarz
(based on relaxation)
- ▶ Physics-split Schur
(based on factorization)
 - ▶ approximate commutators SIMPLE, PCD, LSC
 - ▶ segregated smoothers
 - ▶ Augmented Lagrangian
 - ▶ “parabolization” for stiff waves
- ✗ Need to understand global coupling strengths

- ▶ Preferred data structures depend on which method is used.
- ▶ Interplay with geometric multigrid.

Splitting for Multiphysics

$$\begin{bmatrix} A & B \\ C & D \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} f \\ g \end{bmatrix}$$

- ▶ Relaxation: -pc_fieldsplit_type
[additive,multiplicative,symmetric_multiplicative]

$$\begin{bmatrix} A & D \end{bmatrix}^{-1} \quad \begin{bmatrix} A & \\ C & D \end{bmatrix}^{-1} \quad \begin{bmatrix} A & \\ & 1 \end{bmatrix}^{-1} \left(\mathbf{1} - \begin{bmatrix} A & B \\ & 1 \end{bmatrix} \begin{bmatrix} A & \\ C & D \end{bmatrix}^{-1} \right)$$

- ▶ Gauss-Seidel inspired, works when fields are loosely coupled
- ▶ Factorization: -pc_fieldsplit_type schur

$$\begin{bmatrix} A & B \\ & S \end{bmatrix}^{-1} \begin{bmatrix} 1 & \\ CA^{-1} & 1 \end{bmatrix}^{-1}, \quad S = D - CA^{-1}B$$

- ▶ robust (exact factorization), can often drop lower block
- ▶ how to precondition S which is usually dense?
 - ▶ interpret as differential operators, use approximate commutators

How much nesting?

$$P_1 = \begin{pmatrix} J_{uu} & J_{up} & J_{uE} \\ 0 & B_{pp} & 0 \\ 0 & 0 & J_{EE} \end{pmatrix}$$

$$P = \begin{bmatrix} \begin{pmatrix} J_{uu} & J_{up} \\ J_{pu} & 0 \end{pmatrix} & \\ \begin{pmatrix} J_{Eu} & J_{Ep} \end{pmatrix} & J_{EE} \end{bmatrix}$$

- ▶ B_{pp} is a mass matrix in the pressure space weighted by inverse of kinematic viscosity.
- ▶ Elman, Mihajlović, Wathen, JCP 2011 for non-dimensional isoviscous Boussinesq.
- ▶ Works well for non-dimensional problems on the cube, not for realistic parameters.
 - ▶ Low-order preconditioning full-accuracy unassembled high order operator.
 - ▶ Build these on command line with PETSc PCFieldSplit.
- ▶ Inexact inner solve using upper-triangular with B_{pp} for Schur.
- ▶ Another level of nesting.
- ▶ GCR tolerant of inexact inner solves.
- ▶ Outer converges in 1 or 2 iterations.

Example 3×3 problem with nested 2×2 split

```
-fieldsplit_s_ksp_type gcr
-fieldsplit_s_ksp_rtol 1e-1
-fieldsplit_s_ksp_monitor_vht
-fieldsplit_s_ksp_monitor_singular_value
-fieldsplit_s_pc_type fieldsplit
-fieldsplit_s_pc_fieldsplit_type schur
-fieldsplit_s_pc_fieldsplit_real_diagonal
-fieldsplit_s_pc_fieldsplit_schur_factorization_type lower
-fieldsplit_s_fieldsplit_u_ksp_type gmres
-fieldsplit_s_fieldsplit_u_ksp_max_it 10
-fieldsplit_s_fieldsplit_u_pc_type asm
-fieldsplit_s_fieldsplit_u_sub_pc_type ilu
-fieldsplit_s_fieldsplit_u_sub_pc_factor_levels 1
-fieldsplit_s_fieldsplit_u_ksp_converged_reason
-fieldsplit_s_fieldsplit_p_ksp_type preonly
-fieldsplit_s_fieldsplit_p_ksp_max_it 1
-fieldsplit_s_fieldsplit_p_pc_type jacobi
-fieldsplit_e_ksp_type gmres
-fieldsplit_e_ksp_converged_reason
-fieldsplit_e_pc_type asm
-fieldsplit_e_sub_pc_type ilu
-fieldsplit_e_sub_pc_factor_levels 2
```

Monolithic approaches

Parallel direct solver

```
-dm_mat_type aij -pc_type lu -pc_factor_mat_solver_package mumps
```

Coupled nonlinear multigrid accelerated by NGMRES with multi-stage smoothers

```
-lidvelocity 200 -grashof 1e4  
-snes_grid_sequence 5 -snes_monitor -snes_view  
-snes_type ngmres  
-npc_snes_type fas  
-npc_snes_max_it 1  
-npc_fas_coarse_snes_type ls  
-npc_fas_coarse_ksp_type preonly  
-npc_fas_snes_type ms  
-npc_fas_snes_max_it 1  
-npc_fas_ksp_type preonly  
-npc_fas_pc_type pbjacobi  
-npc_fas_snes_ms_type m62  
-npc_fas_snes_max_it 1
```

Outline

Multiphysics and methods

Coupling software in PETSc

Hardware and consequences

Multi-physics coupling in PETSc

Momentum

Pressure

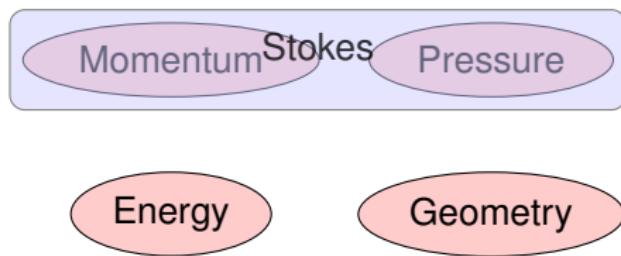
- ▶ package each “physics” independently
- ▶ solve single-physics and coupled problems
- ▶ semi-implicit and fully implicit
- ▶ reuse residual and Jacobian evaluation unmodified
- ▶ direct solvers, fieldsplit inside multigrid, multigrid inside fieldsplit without recompilation
- ▶ use the best possible matrix format for each physics (e.g. symmetric block size 3)
- ▶ matrix-free anywhere
- ▶ multiple levels of nesting

Multi-physics coupling in PETSc

Momentum Stokes Pressure

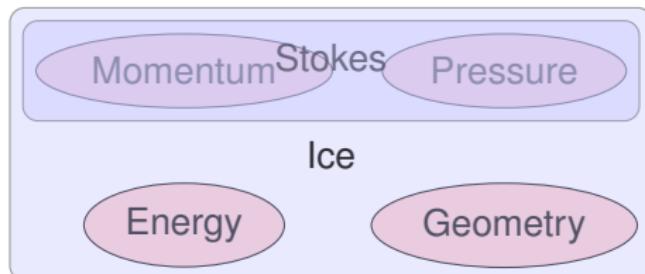
- ▶ package each “physics” independently
- ▶ solve single-physics and coupled problems
- ▶ semi-implicit and fully implicit
- ▶ reuse residual and Jacobian evaluation unmodified
- ▶ direct solvers, fieldsplit inside multigrid, multigrid inside fieldsplit without recompilation
- ▶ use the best possible matrix format for each physics (e.g. symmetric block size 3)
- ▶ matrix-free anywhere
- ▶ multiple levels of nesting

Multi-physics coupling in PETSc



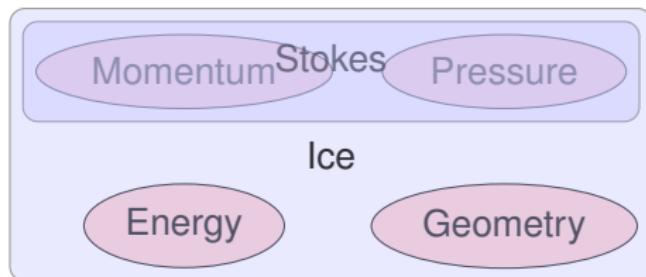
- ▶ package each “physics” independently
- ▶ solve single-physics and coupled problems
- ▶ semi-implicit and fully implicit
- ▶ reuse residual and Jacobian evaluation unmodified
- ▶ direct solvers, fieldsplit inside multigrid, multigrid inside fieldsplit without recompilation
- ▶ use the best possible matrix format for each physics (e.g. symmetric block size 3)
- ▶ matrix-free anywhere
- ▶ multiple levels of nesting

Multi-physics coupling in PETSc

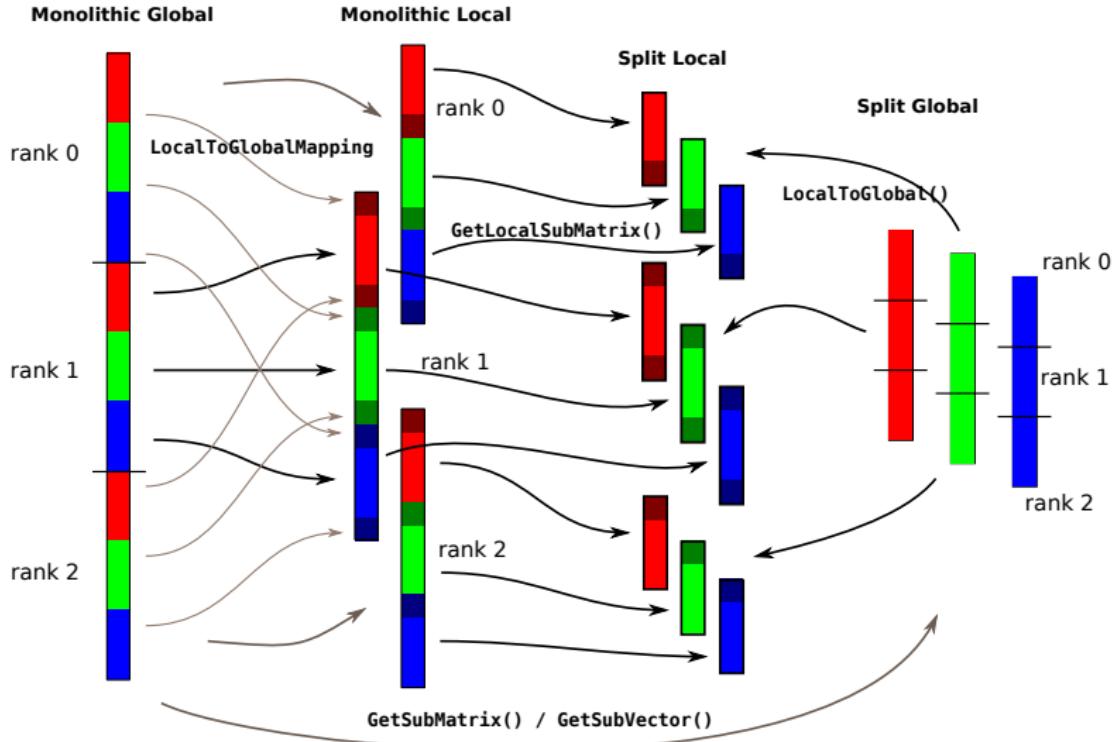


- ▶ package each “physics” independently
- ▶ solve single-physics and coupled problems
- ▶ semi-implicit and fully implicit
- ▶ reuse residual and Jacobian evaluation unmodified
- ▶ direct solvers, fieldsplit inside multigrid, multigrid inside fieldsplit without recompilation
- ▶ use the best possible matrix format for each physics (e.g. symmetric block size 3)
- ▶ matrix-free anywhere
- ▶ multiple levels of nesting

Multi-physics coupling in PETSc



- ▶ package each “physics” independently
- ▶ solve single-physics and coupled problems
- ▶ semi-implicit and fully implicit
- ▶ reuse residual and Jacobian evaluation unmodified
- ▶ direct solvers, fieldsplit inside multigrid, multigrid inside fieldsplit without recompilation
- ▶ use the best possible matrix format for each physics (e.g. symmetric block size 3)
- ▶ matrix-free anywhere
- ▶ multiple levels of nesting



Work in Split Local space, matrix data structures reside in any space.

```
MatGetLocalSubMatrix(Mat A,IS rows,IS cols,Mat *B);
```

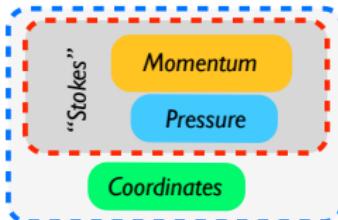
- ▶ Primarily for assembly
 - ▶ B is not guaranteed to implement MatMult
 - ▶ The communicator for B is not specified,
only safe to use non-collective ops (unless you check)
- ▶ IS represents an index set, includes a block size and
communicator
- ▶ MatSetValuesBlockedLocal() is implemented
- ▶ MatNest returns nested submatrix, no-copy
- ▶ No-copy for Neumann-Neumann formats
(unassembled across procs, e.g. BDDC, FETI-DP)
- ▶ Most other matrices return a lightweight proxy Mat
 - ▶ COMM_SELF
 - ▶ Values not copied, does not implement MatMult
 - ▶ Translates indices to the language of the parent matrix
 - ▶ Multiple levels of nesting are flattened

Stokes + Implicit Free Surface

$$[\eta D_{ij}(\mathbf{u})]_{,j} - p_{,i} = f_i$$

$$u_{k,k} = 0$$

$$\hat{x}_i = \hat{x}_i^{t-\Delta t} + \Delta t u_i(\hat{x}_i)$$



COORDINATE RESIDUALS

$$F_x := -u_i + \frac{\hat{x}_i}{\Delta t} - \frac{\hat{x}_i^{t-\Delta t}}{\Delta t}$$

[We use a full Lagrangian update of our mesh, with no remeshing]

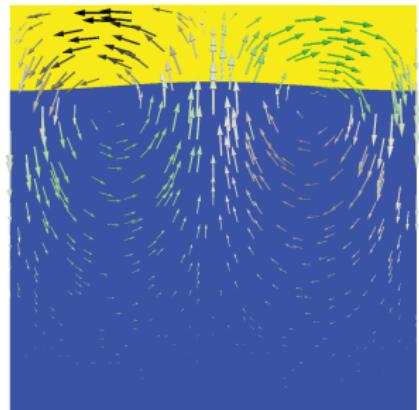
JACOBIAN

$$\mathcal{J}_{si} = \begin{bmatrix} A + \delta_{\hat{x}} A & B + \delta_{\hat{x}} B \\ B^T + \delta_{\hat{x}} B^T & 0 \\ -I & 0 \end{bmatrix} \begin{bmatrix} J_{ac} \\ J_{bc} \\ \frac{I}{\Delta t} \end{bmatrix}$$

Reuse stokes operators and saddle point preconditioners

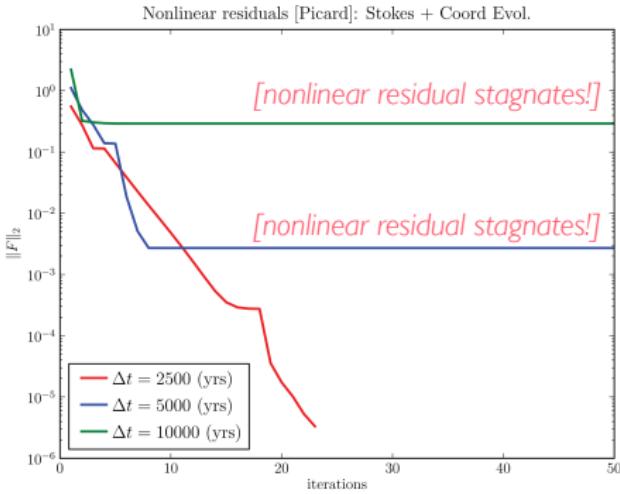
NESTED PRECONDITIONER

$$\mathcal{P}_{si} = \left[\begin{bmatrix} \mathcal{P}_s^l \\ I \end{bmatrix} \left[\begin{bmatrix} -I \\ -\frac{I}{\Delta t} \end{bmatrix} \right] \right] \mathcal{P}_s^l = \begin{bmatrix} A & 0 \\ B^T & -S \end{bmatrix}$$

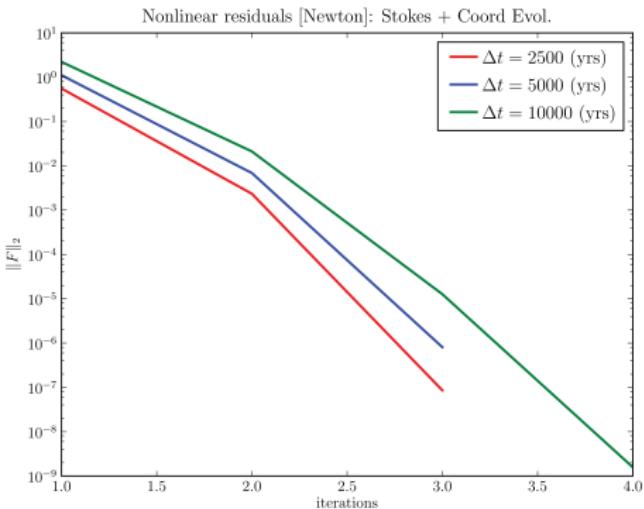


"Drunken seaman", Rayleigh Taylor instability test case from Kaus et al., 2010. Dense, viscous material (yellow) overlying less dense, less viscous material (blue).

Stokes + Implicit Free Surface

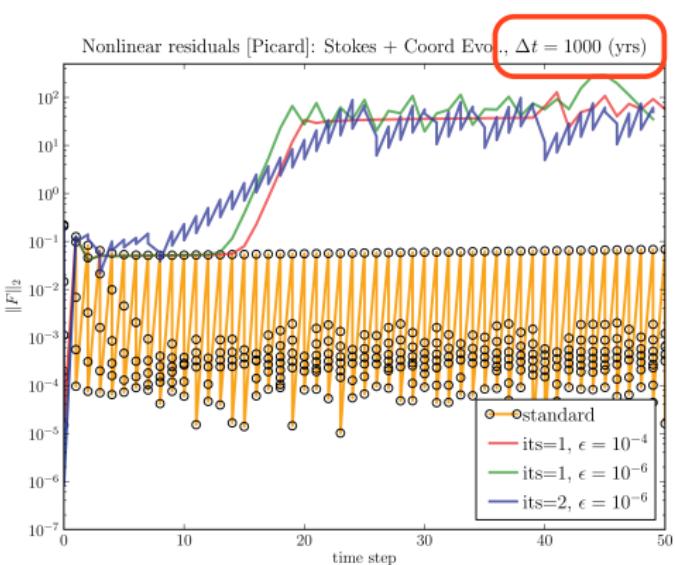
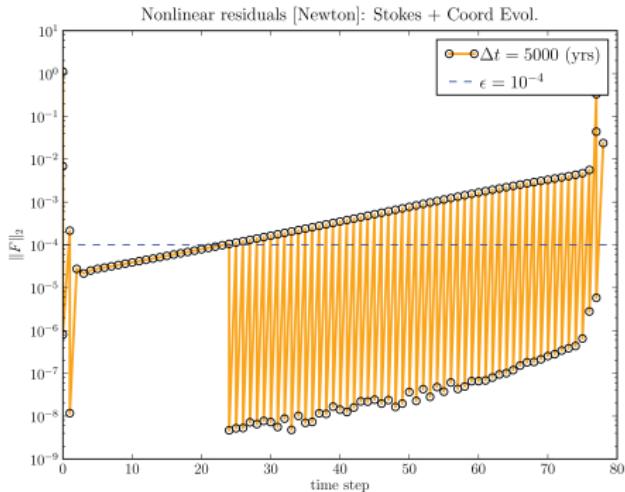


* Picard fails to converge for large time step sizes.



* Newton is robust for a wide range of time step sizes.

Stokes + Implicit Free Surface



- * The nonlinear residual **ALWAYS** increases from one step to the next.
- * A nonlinear solve is required to control the error.

- * An accurate nonlinear solve on the first time step, combined with 1 or 2 nonlinear iterations on subsequent steps still results in severe errors. *This is true even when dt is small.*

Outline

Multiphysics and methods

Coupling software in PETSc

Hardware and consequences

On-node hardware roadmap

Hardware trends

- ▶ More cores (keep hearing $\mathcal{O}(1000)$ per node)
- ▶ Long vector registers (32B for AVX and BG/Q, 64B for MIC)
- ▶ Must use SMT to hide memory latency (POWER7)
- ▶ Must use SMT for floating point performance (GPU, BG/Q)
- ▶ Large penalty for non-contiguous memory access

“Free flops”, but how can we use them?

- ▶ High order methods good: better accuracy per storage
- ▶ High order methods bad: work unit gets larger
- ▶ GPU threads have very little memory, must keep work unit small
- ▶ Need library composability, keep user contribution embarrassingly parallel

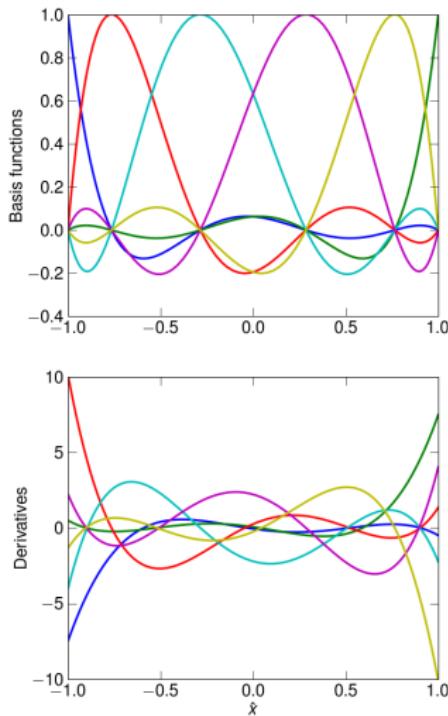
How to program this beast?

- ▶ Decouple physics from discretization
 - ▶ Expose small, embarrassingly parallel operations to user
 - ▶ Library schedules user threads for reuse between kernels
 - ▶ User provides physics in kernels run at each quadrature point
 - ▶ Continuous weak form: find $u \in \mathcal{V}_D$

$$v^T F(u) \sim \int_{\Omega} v \cdot f_0(u, \nabla u) + \nabla v : f_1(u, \nabla u) = 0, \quad \forall v \in \mathcal{V}_0$$

- ▶ Similar form at faces, may involve Riemann solve
- ▶ Library manages reductions
 - ▶ Interpolation and differentiation on elements
 - ▶ Exploit tensor product structure to keep working set small
 - ▶ Assembly into solution/residual vector (sum over elements)

Nodal hp -version finite element methods



1D reference element

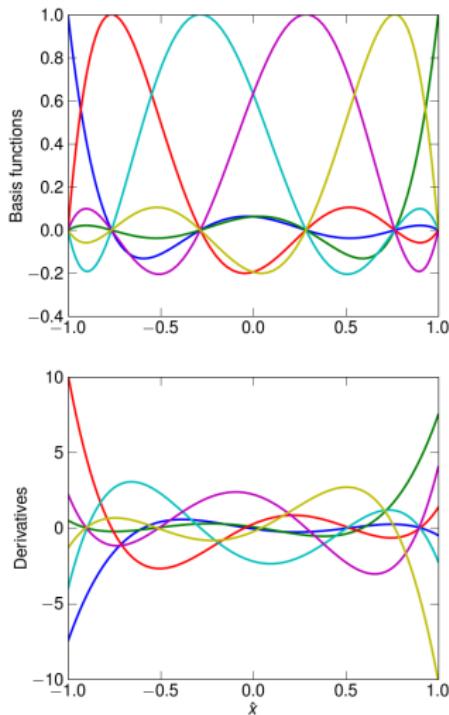
- ▶ Lagrange interpolants on Legendre-Gauss-Lobatto points
- ▶ Quadrature \hat{R} , weights \hat{W}
- ▶ Evaluation: \hat{B}, \hat{D}

3D reference element

$$\begin{aligned}\hat{W} &= \hat{W} \otimes \hat{W} \otimes \hat{W} \\ \hat{B} &= \hat{B} \otimes \hat{B} \otimes \hat{B} \\ \hat{D}_0 &= \hat{D} \otimes \hat{B} \otimes \hat{B} \\ \hat{D}_1 &= \hat{B} \otimes \hat{D} \otimes \hat{B} \\ \hat{D}_2 &= \hat{B} \otimes \hat{B} \otimes \hat{D}\end{aligned}$$

These tensor product operations
are very efficient, 70% of peak flop/s

Nodal hp -version finite element methods



1D reference element

- ▶ Lagrange interpolants on Legendre-Gauss-Lobatto points
- ▶ Quadrature \hat{R} , weights \hat{W}
- ▶ Evaluation: \hat{B}, \hat{D}

3D reference element

$$\begin{aligned}\hat{W} &= \hat{W} \otimes \hat{W} \otimes \hat{W} \\ \hat{B} &= \hat{B} \otimes \hat{B} \otimes \hat{B} \\ \hat{D}_0 &= \hat{D} \otimes \hat{B} \otimes \hat{B} \\ \hat{D}_1 &= \hat{B} \otimes \hat{D} \otimes \hat{B} \\ \hat{D}_2 &= \hat{B} \otimes \hat{B} \otimes \hat{D}\end{aligned}$$

These tensor product operations
are very efficient, 70% of peak flop/s

Operations on physical elements

Mapping to physical space

$$x^e : \hat{K} \rightarrow K^e, \quad J_{ij}^e = \partial x_i^e / \partial \hat{x}_j, \quad (J^e)^{-1} = \partial \hat{x} / \partial x^e$$

Element operations in physical space

$$B^e = \hat{B} \quad W^e = \hat{W} \Lambda(|J^e(r)|)$$

$$D_i^e = \Lambda\left(\frac{\partial \hat{x}_0}{\partial x_i}\right) \hat{D}_0 + \Lambda\left(\frac{\partial \hat{x}_1}{\partial x_i}\right) \hat{D}_1 + \Lambda\left(\frac{\partial \hat{x}_2}{\partial x_i}\right) \hat{D}_2$$

$$(D_i^e)^T = \hat{D}_0^T \Lambda\left(\frac{\partial \hat{x}_0}{\partial x_i}\right) + \hat{D}_1^T \Lambda\left(\frac{\partial \hat{x}_1}{\partial x_i}\right) + \hat{D}_2^T \Lambda\left(\frac{\partial \hat{x}_2}{\partial x_i}\right)$$

Global problem is defined by assembly

$$F(u) = \sum_e \mathcal{E}_e^T \left[(B^e)^T W^e \Lambda(\mathbf{f}_0(u^e, \nabla u^e)) + \sum_{i=0}^d (D_i^e)^T W^e \Lambda(f_{1,i}(u^e, \nabla u^e)) \right] = 0$$

where $u^e = B^e \mathcal{E}^e u$ and $\nabla u^e = \{D_i^e \mathcal{E}^e u\}_{i=0}^2$

Representation of Jacobians, Automation

- ▶ For unassembled representations, decomposition, and assembly
- ▶ Continuous weak form: find u

$$v^T F(u) \sim \int_{\Omega} v \cdot f_0(u, \nabla u) + \nabla v : f_1(u, \nabla u) = 0, \quad \forall v \in \mathcal{V}_0$$

- ▶ Weak form of the Jacobian $J(u)$: find w

$$v^T J(u) w \sim \int_{\Omega} [v^T \quad \nabla v^T] \begin{bmatrix} f_{0,0} & f_{0,1} \\ f_{1,0} & f_{1,1} \end{bmatrix} \begin{bmatrix} w \\ \nabla w \end{bmatrix}$$
$$[f_{i,j}] = \begin{bmatrix} \frac{\partial f_0}{\partial u} & \frac{\partial f_0}{\partial \nabla u} \\ \frac{\partial f_1}{\partial u} & \frac{\partial f_1}{\partial \nabla u} \end{bmatrix} (u, \nabla u)$$

- ▶ Terms in $[f_{i,j}]$ easy to compute symbolically, AD more scalable.
- ▶ Nonlinear terms f_0, f_1 usually have the most expensive nonlinearities in the computation of scalar material parameters
 - ▶ Equations of state, effective viscosity, “star” region in Riemann solve
 - ▶ Compute gradient with reverse-mode, store at quadrature points.
 - ▶ Perturb scalars, then use forward-mode to complete the Jacobian.
 - ▶ Flip for action of the adjoint.

Conservative (non-Boussinesq) two-phase ice flow

Find momentum density ρu , pressure p , and total energy density E :

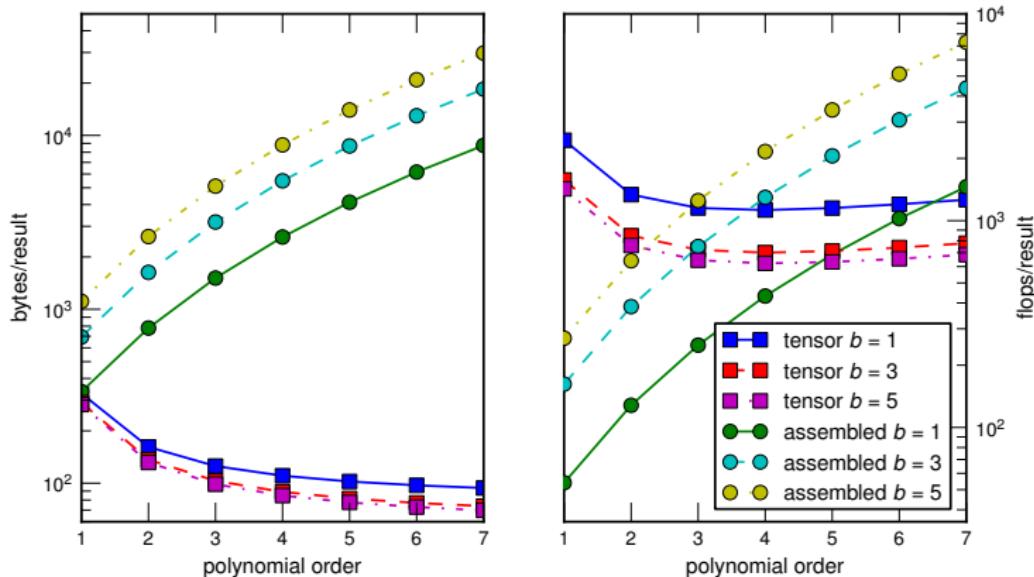
$$(\rho u)_t + \nabla \cdot (\rho u \otimes u - \eta D u_i + p \mathbf{1}) - \rho g = 0$$

$$\rho_t + \nabla \cdot \rho u = 0$$

$$E_t + \nabla \cdot ((E + p)u - k_T \nabla T - k_\omega \nabla \omega) - \eta D u_i : D u_i - \rho u \cdot g = 0$$

- ▶ Solve for density ρ , ice velocity u_i , temperature T , and melt fraction ω using constitutive relations.
 - ▶ Simplified constitutive relations can be solved explicitly.
 - ▶ Temperature, moisture, and strain-rate dependent rheology η .
 - ▶ High order FEM, typically Q_3 momentum & energy
- ▶ DAEs solved implicitly after semidiscretizing in space.
- ▶ Preconditioning using nested fieldsplit

Performance of assembled versus unassembled



- ▶ High order Jacobian stored unassembled using coefficients at quadrature points, can use local AD
- ▶ Choose approximation order at run-time, independent for each field
- ▶ Precondition high order using assembled lowest order method
- ▶ Implementation > 70% of FPU peak, SpMV bandwidth wall < 4%

Memory Bandwidth

Operation	Arithmetic Intensity (flops per byte)		
Sparse matrix-vector product	1/6		
Dense matrix-vector product	1/4		
Unassembled matrix-vector product	≈ 8		
High-order residual evaluation	> 5		
Processor	BW (GB/s)	Peak (GF/s)	Balanced AI (F/B)
Sandy Bridge 6-core	21*	150	7.2
Magny Cours 16-core	42*	281	6.7
Blue Gene/Q node	43	205	4.8
GeForce 9400M	21	54	2.6
GTX 285	159	1062	6.8
Tesla M2050	144	1030	7.1

Prospects for reducing synchronization

- ▶ Dot products and norms
 - ▶ orthogonality is a powerful concept
 - ▶ dot product/norm fusion in CG variants
 - ▶ latency-tolerant Krylov methods, TSQR for GMRES
 - ▶ nonlinear methods (e.g. NGMRES, BFGS, line searches)
 - ▶ hierarchical methods to limit system-wide norms
 - ▶ setting up smoothers and coarsening rates for AMG
- ▶ additive coarse grids
- ▶ subphysics on subcommunicators, even within multigrid context
- ▶ s -step methods (and other fusion)
 - ▶ often spoiled by algorithmic requirements of preconditioning
 - ▶ relevant for multigrid smoothers
 - ▶ difficult crossovers for 3D problems

Multigrid is *always* strong scaling

- ▶ Finest level is chosen by the application (might have big subdomains)
- ▶ All coarsened levels choose communicator size based on strong scaling limit
- ▶ Optimizing the strong scaling limit pays off consistently
- ▶ Rapid coarsening is important (2:1 semi-coarsening not okay any more)

Advanced Analysis

- ▶ Uncertainty quantification
 - ▶ intrusive vs unintrusive methods, multilevel
 - ▶ uncertainty in modeling error
 - ▶ use subdifferentials for non-smooth processes
 - ▶ unified handling of heterogeneous observational data
- ▶ PDE-constrained optimization
 - ▶ multi-objective
 - ▶ robustness
 - ▶ rich problem description
 - ▶ fusing algorithmic steps (LNK and coupled DD fuse gradients with progress)
- ▶ Exploring stability manifolds
 - ▶ solving bordered linear and nonlinear systems
- ▶ nearly time-periodic nonlinear problems
 - ▶ identifying cycles in ocean models, turbomachinery

Software challenges

- ▶ Which interfaces do users have to interact with?
 - ▶ “F”ramework vs library
 - ▶ Extensibility is critical for multiphysics
- ▶ Asynchronous interfaces crossing module boundaries
 - ▶ How to ensure progress?
- ▶ Merge communication on multiple levels or between multiple physics
- ▶ Fusing coarse level operations
- ▶ Working with non-nested communicators is tricky
- ▶ Current solutions for hierarchical memory are bad for libraries
 - ▶ I want a communicator-like object
 - ▶ I want a way to allocate memory explicitly/relative to algorithmic dependencies instead of implicit “first touch”
- ▶ Time integration: IMEX, multirate, parallel in time
 - ▶ method of lines: $g(\dot{u}, u, t) = f(u, t)$
 - ▶ Lax-Wendroff time integration is harder for composable software