



IBM Systems Research

Computational Steering of HPC Applications on Blue Gene Systems

Bruce D'Amora

IBM T.J. Watson Research Center

KAUST Collaborators:

Ying Qian, Madhu Srinivasan, Dan Lecocq, Iain Georgeson, Bob Danani

Outline

- **Motivation**
- **A Framework**
- **RealityGrid**
- **Blue Gene Implementation**
- **Next Steps**

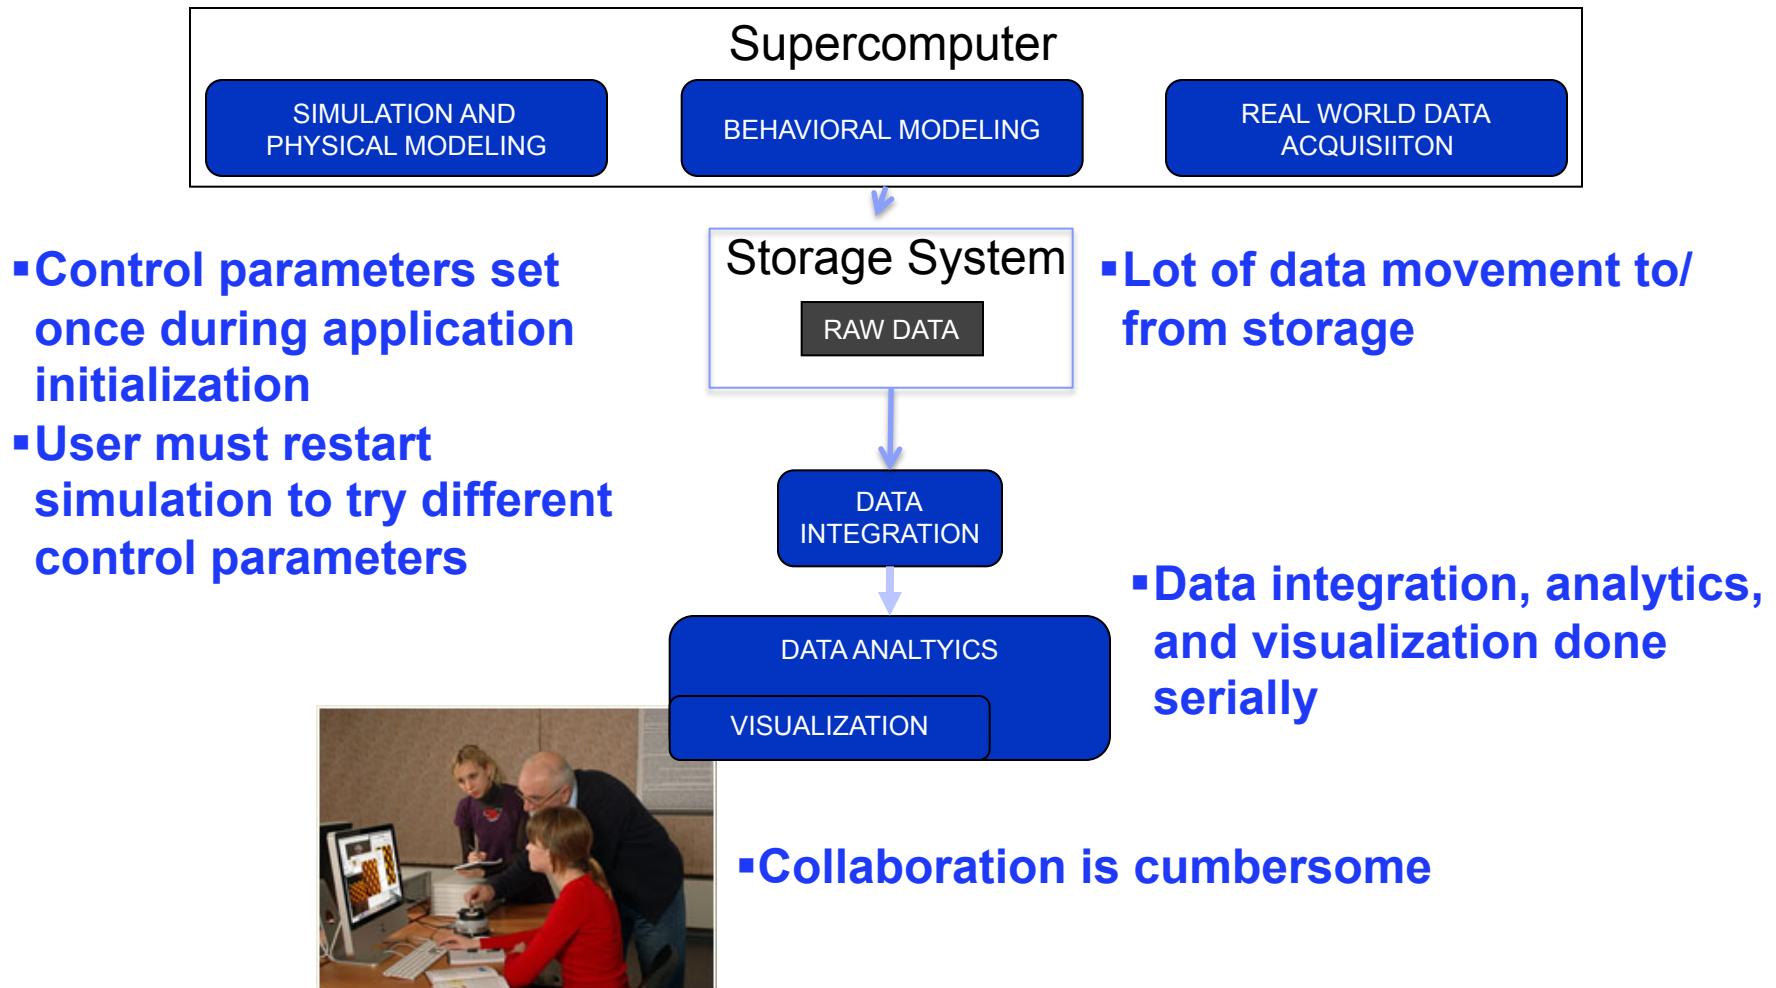
Motivation for Steering

- **Realtime**

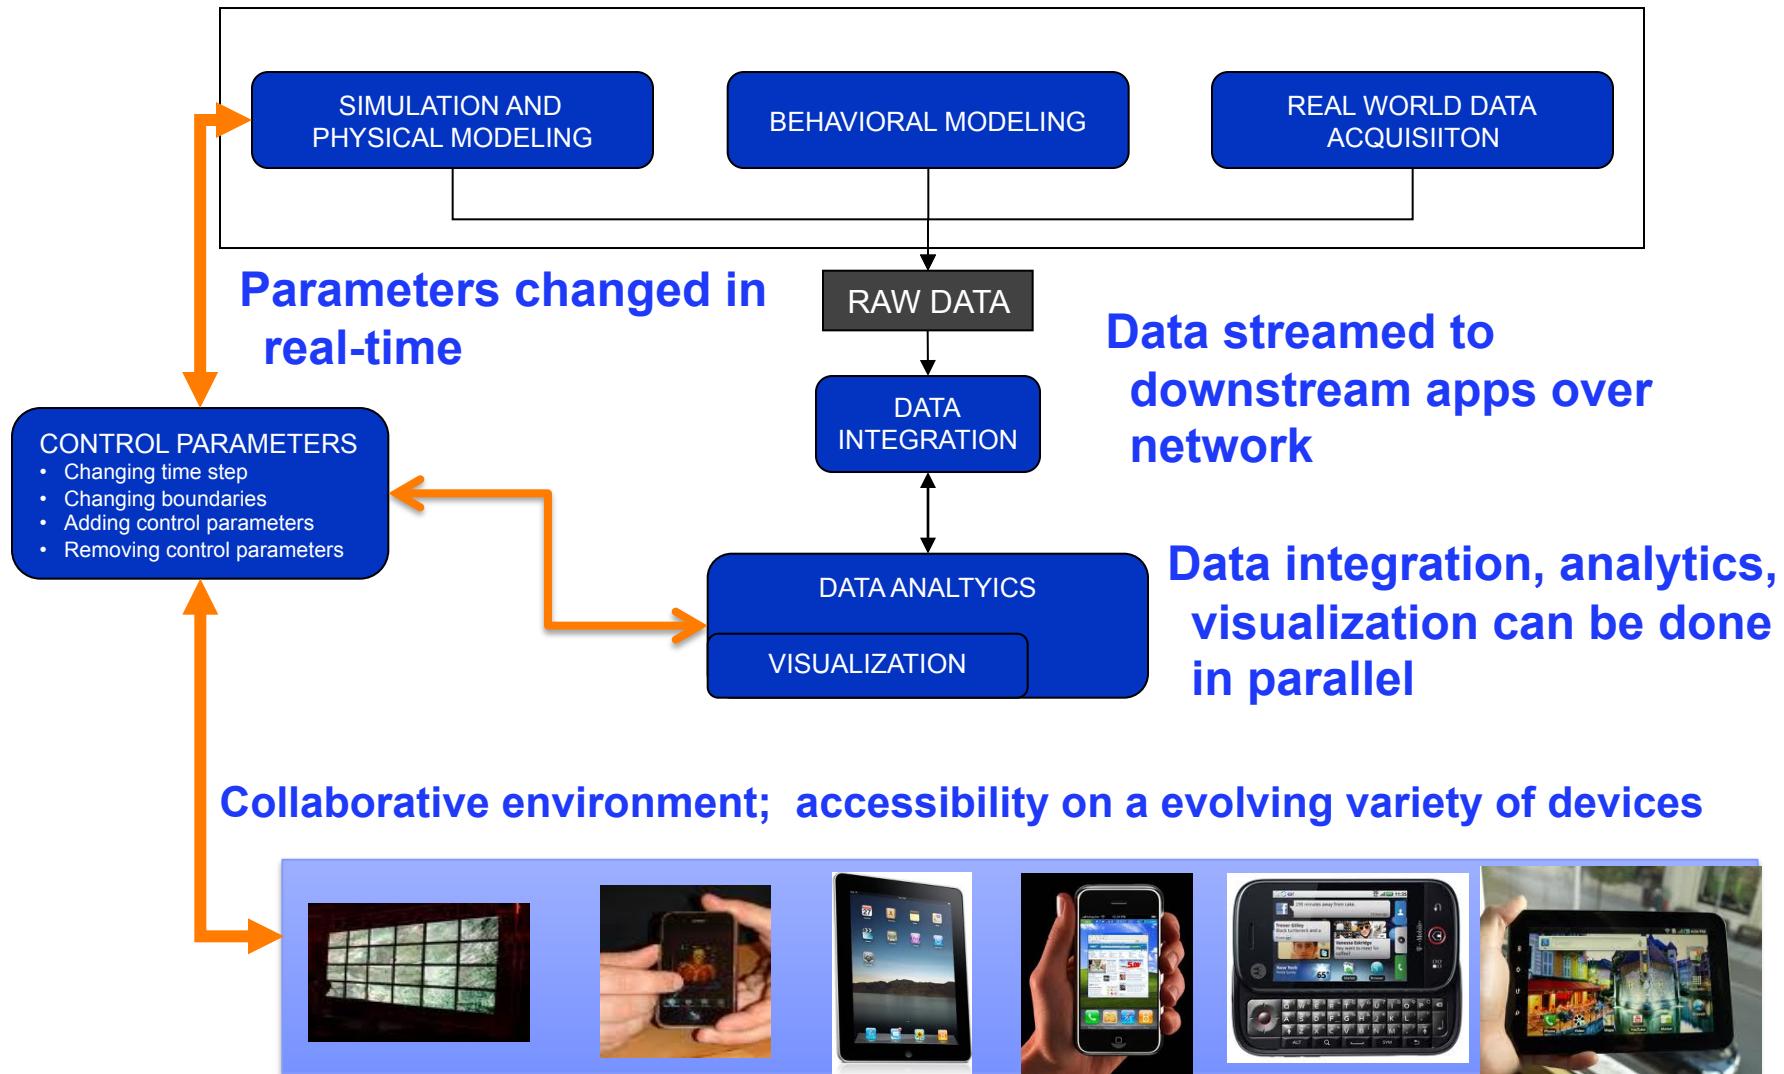
- Coupling of simulation results with visualization or other downstream apps (analytics)
- Exploration of results
- Modification of control parameters

Human Feedback in Simulation Loop

Present Computational Workflow



Future Computational Workflow



Goals for Steering Framework

Usage Scenarios

- User modifiable control parameters during simulation
- Visualization of intermediate & final results of simulations
- Support collaboration – multiple users at multiple locations

Programmability

- Framework must be “thin”; performance sensitive
- Easy to use; complex APIs delay or completely inhibit adoption
- Easily Extensible – no single API fits every applications needs
- Web services enabled – needs to be accessible in the Cloud
- Open source

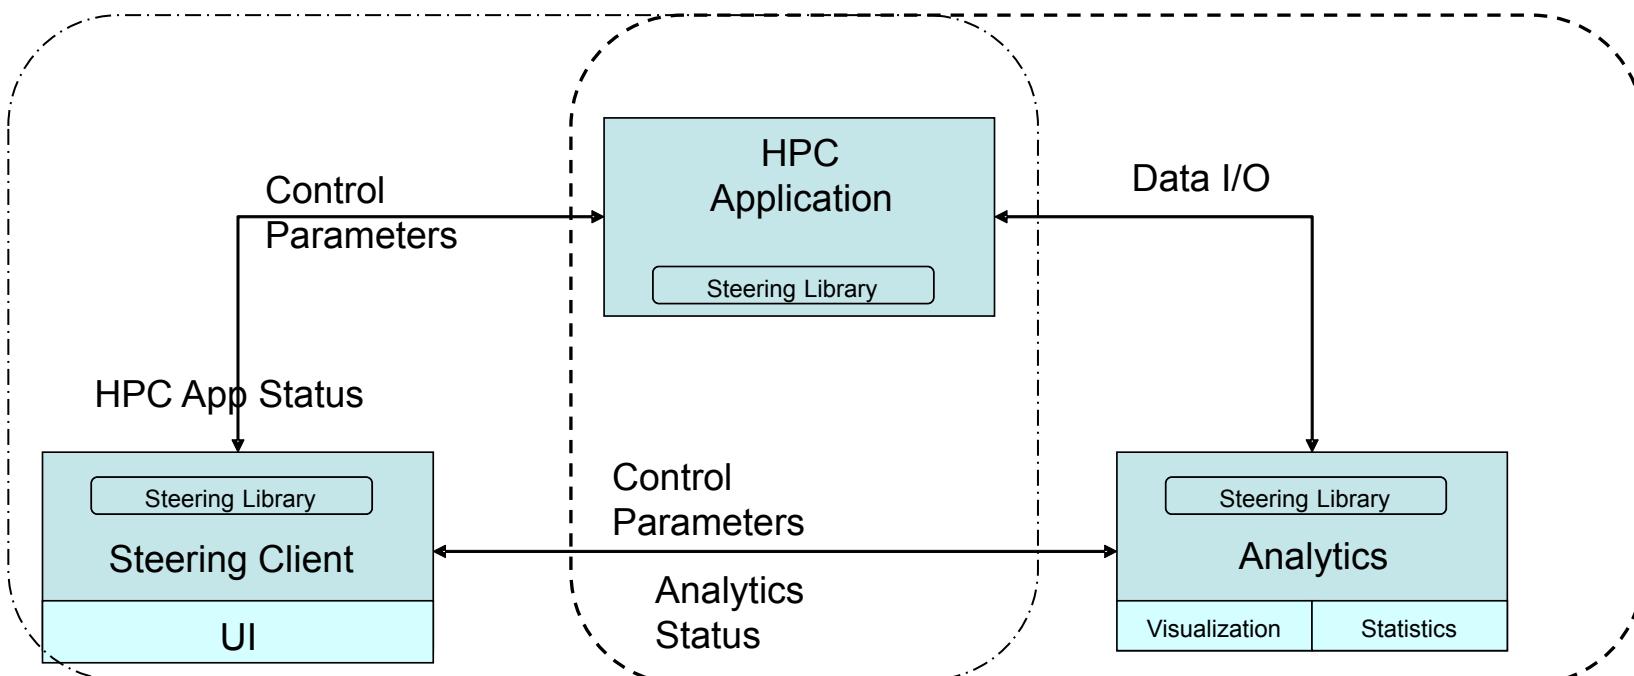
RealityGrid Steering Library

- Developed with UK Engineering and Physical Sciences Research Council (EPSRC) grant.
- Consortium of universities produced BSD licensed Steering Library
 - [University College, London](#) ([Centre for Computational Science](#) and [Department of Chemistry](#))
 - [University of Edinburgh](#) ([Department of Physics and Astronomy](#) and [Edinburgh Parallel Computing Centre](#))
 - [University of Manchester](#) ([Manchester Computing](#) and [Department of Computer Science](#))
 - [Imperial College, University of London](#) ([Department of Computing](#))
 - [Loughborough University](#) ([Loughborough University Advanced VR Research Centre](#) and [Department of Mathematical Sciences](#))
 - [University of Oxford](#) ([Department of Materials](#)).
- The collaborating organizations include the
 - [Computation for Science Consortium](#) (which operates the UK's national supercomputing facility [CSAR](#) in Manchester)
 - [Schlumberger Cambridge Research Ltd.](#)
 - [Edward Jenner Institute for Vaccine Research](#)
 - [Silicon Graphics Inc.](#)
 - [Advanced Visual Systems Inc.](#)
 - [Fujitsu Ltd.](#)
 - [Btexact](#) (in collaboration with [University of Manchester](#))
- Principal Investigator: Peter Coveney, University College London

RealityGrid Features

- **Supports 3 types of communications between HPC application and steering clients**
 - Direct socket connection
 - File based communications – closest to current workflow
 - Web Services Resource Framework (WSRF-Lite) interfaces
- **All communications are embedded in xml streams**
- **WSRF provides persistent state and capability of registering multiple HPC applications to the grid framework**
 - Clients can pick which HPC applications they wish to connect with via web interface
- **Includes bindings for C/C++, Fortran, python, perl, java, wrappers**
- **VTK data reader class developed**

RealityGrid Steering Framework



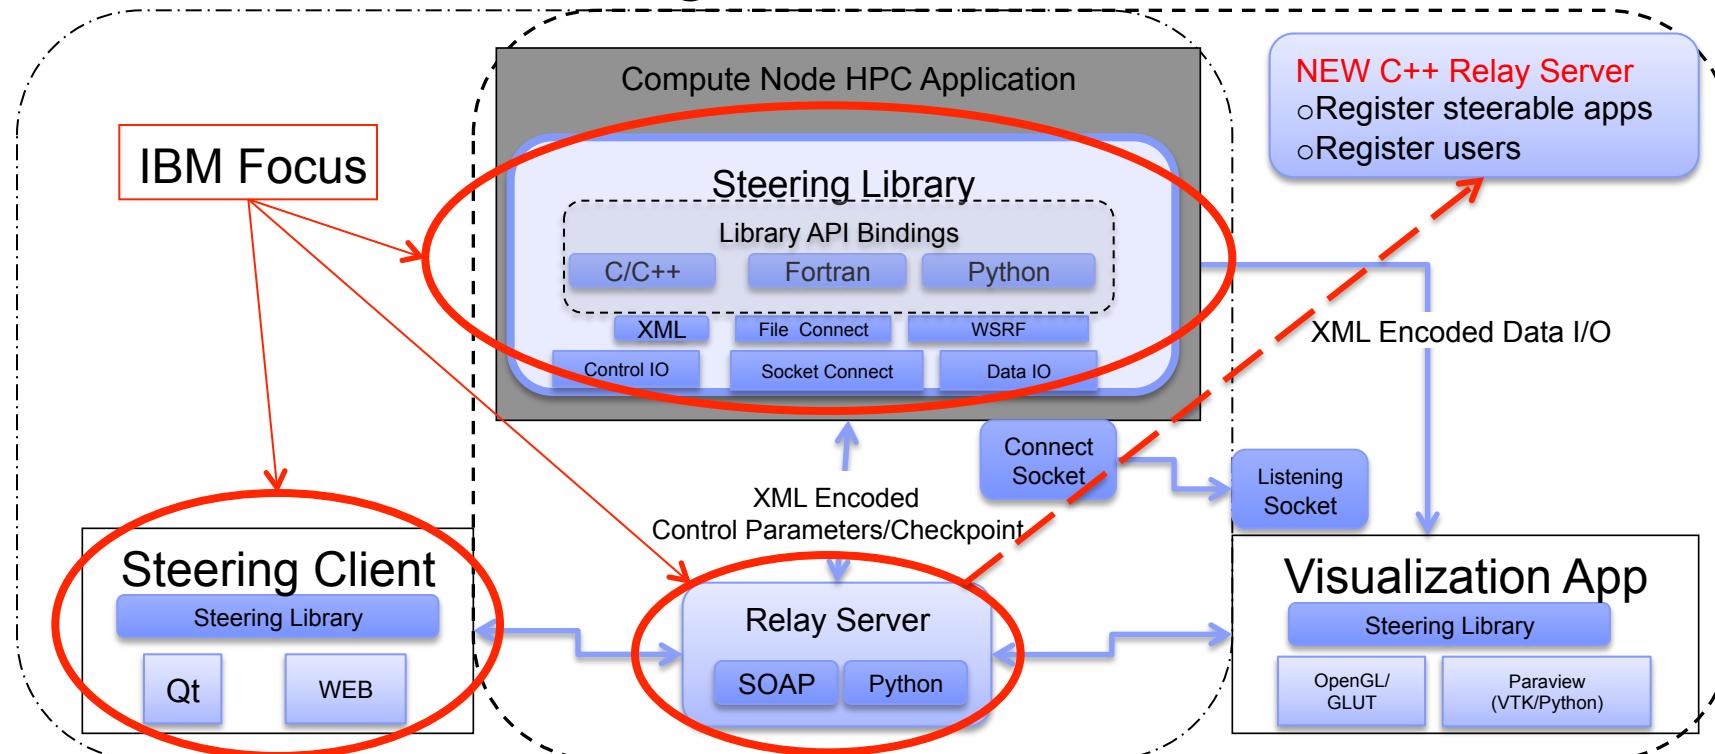
- **Steering Library Function** include, but are not limited to:
 - Identify connections between HPC, Steering, Analytics
 - Identify control parameters as read, write, update
 - Identify data as read, write, update, float, char, integer, etc.
 - Support multiple connections from different clients
 - Query status of connections and applications
 - Encode commands & parameters in XML stream



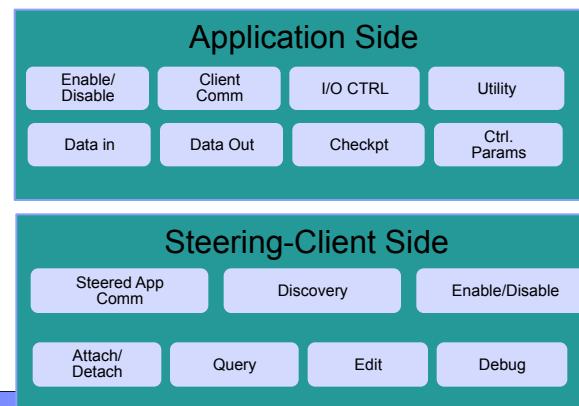
IBM Blue Gene Implementation

- **API implicitly couples IO direction with type of socket connection**
 - i.,e. Output data coincided with a listening socket; Input data coincided with connecting socket
 - Problematic for BlueGene because of limitations of CNK
 - CN (ION) can only establish connecting sockets
- **Solution was to create new API for specifying I/O**
 - Decouple data flow from socket type
- **BlueGene front-end node behind corporate or university firewalls**
 - No direct access to Compute or IO nodes
 - Developed a relay server to allow interaction across firewalls
 - Server also keeps track of steerable applications and the clients that are currently connected
- **Bindings for Python, Perl, and Java**
- **VTK wrappers to abstract the RealityGrid steering interface into a “data reader” class have also been written**
- **Enhanced steering clients: web or native application**

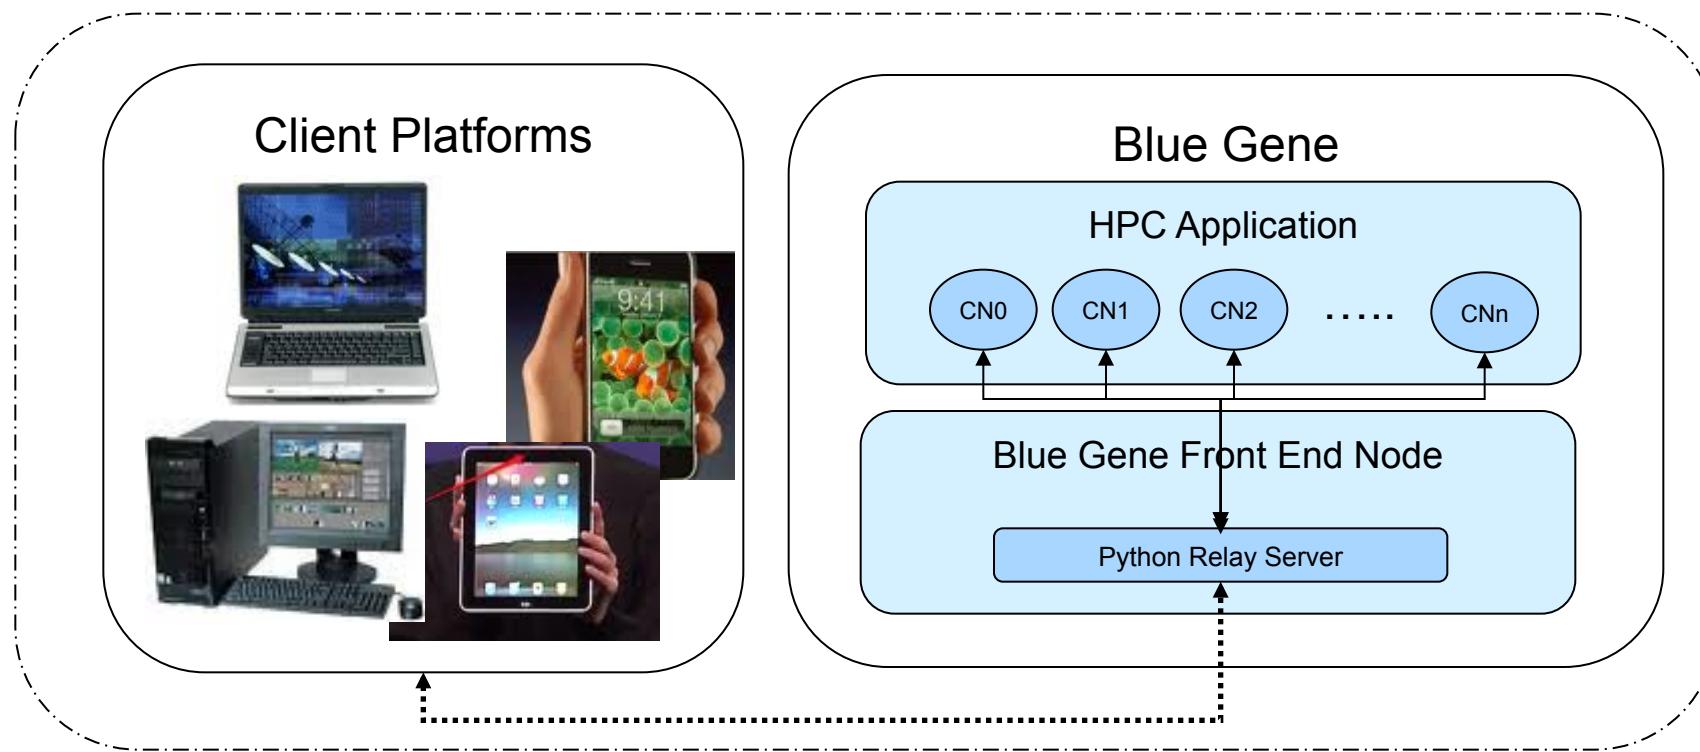
Blue Gene Steering Framework



- **Steering Library Functions** include, but are not limited to:
 - Establish connections between HPC, Steering, Analytics applications running on different platforms across firewalls
 - Register control parameters as read, write, update and transfer values between enabled applications
 - Register data as read, write, update, float, char, integer, etc. and transfer between enabled applications
 - Support multiple connections from different clients
 - Record and transfer checkpoints
 - Query status of connections and applications
 - Encode commands & parameters in XML stream to easily transfer over http connections



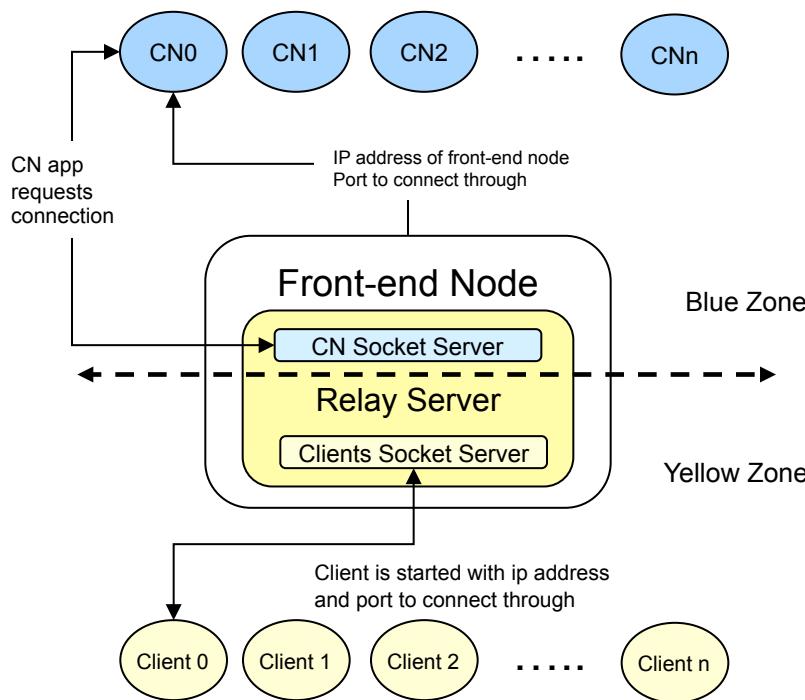
Command and Parameter Flow



- Python Relay Program: Socket Client-Server interface
 - Allows CN to communicate with “outside” world
 - Relay client-side: Establishes in-bound connection with CN socket
 - Relay server-side: Binds connection requests from multiple steering clients
 - Pass thru commands and parameters

Relay Server

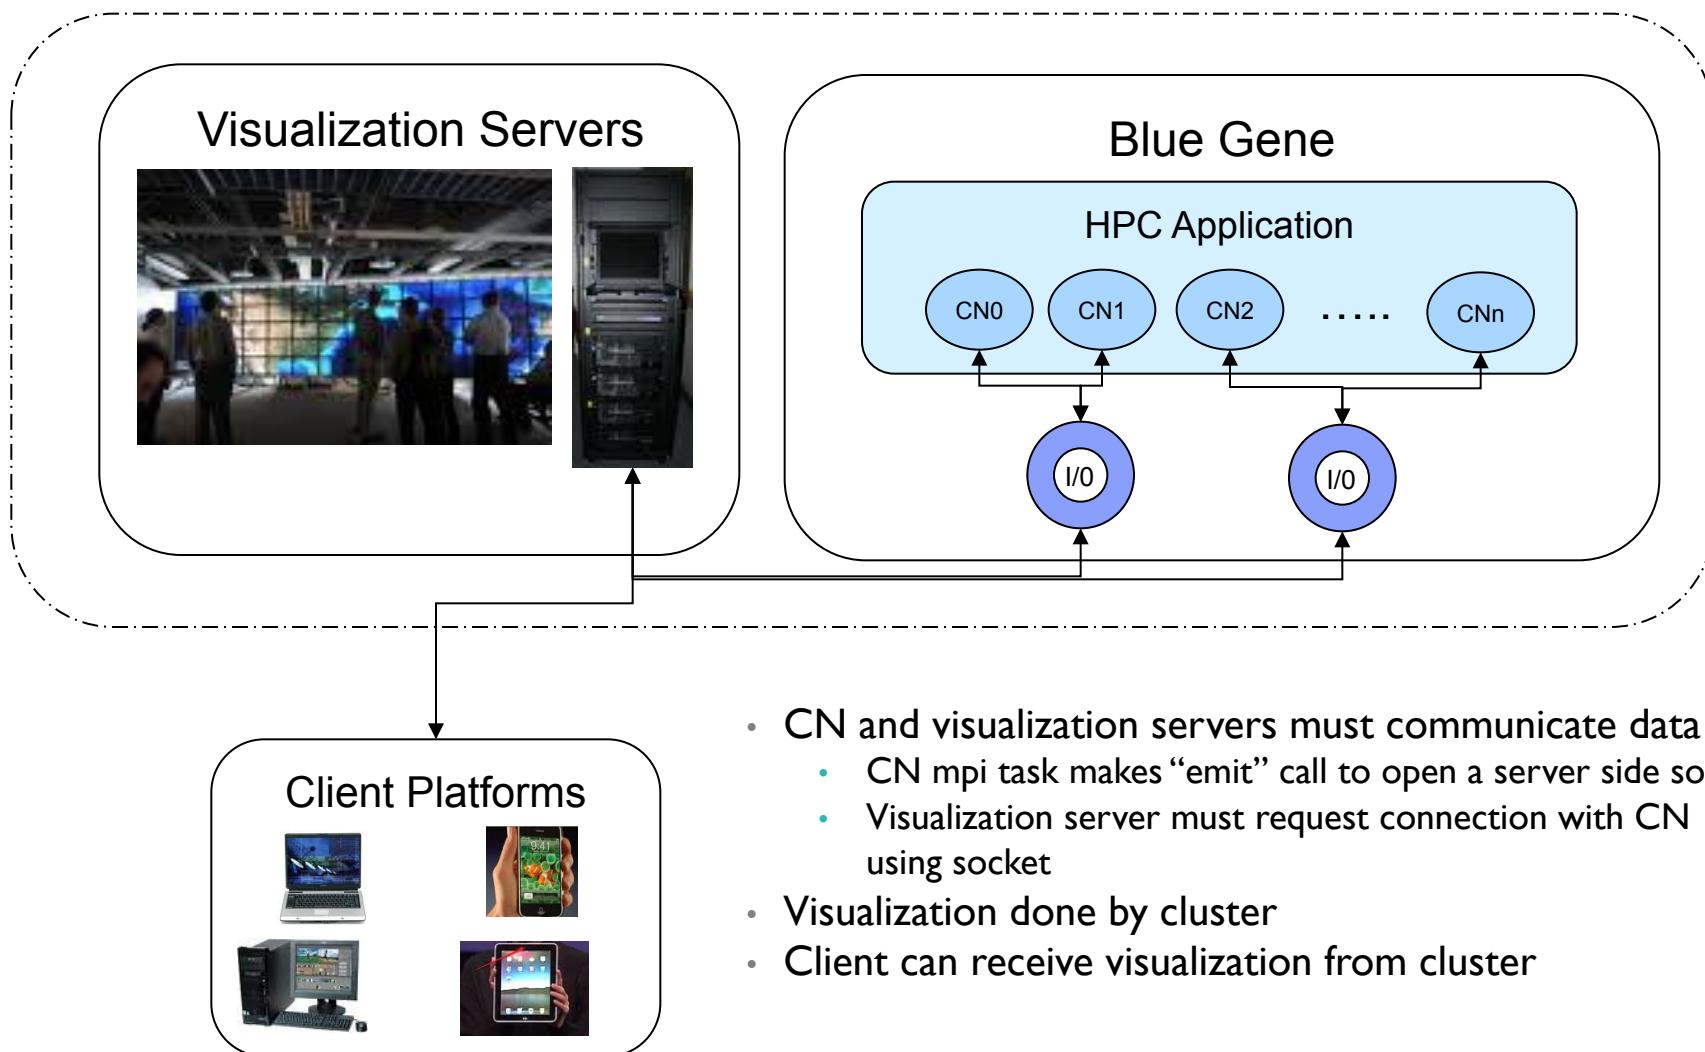
- Compute Nodes on private network so inbound connections are problematic
- One solution is to run relay socket-client on FEN



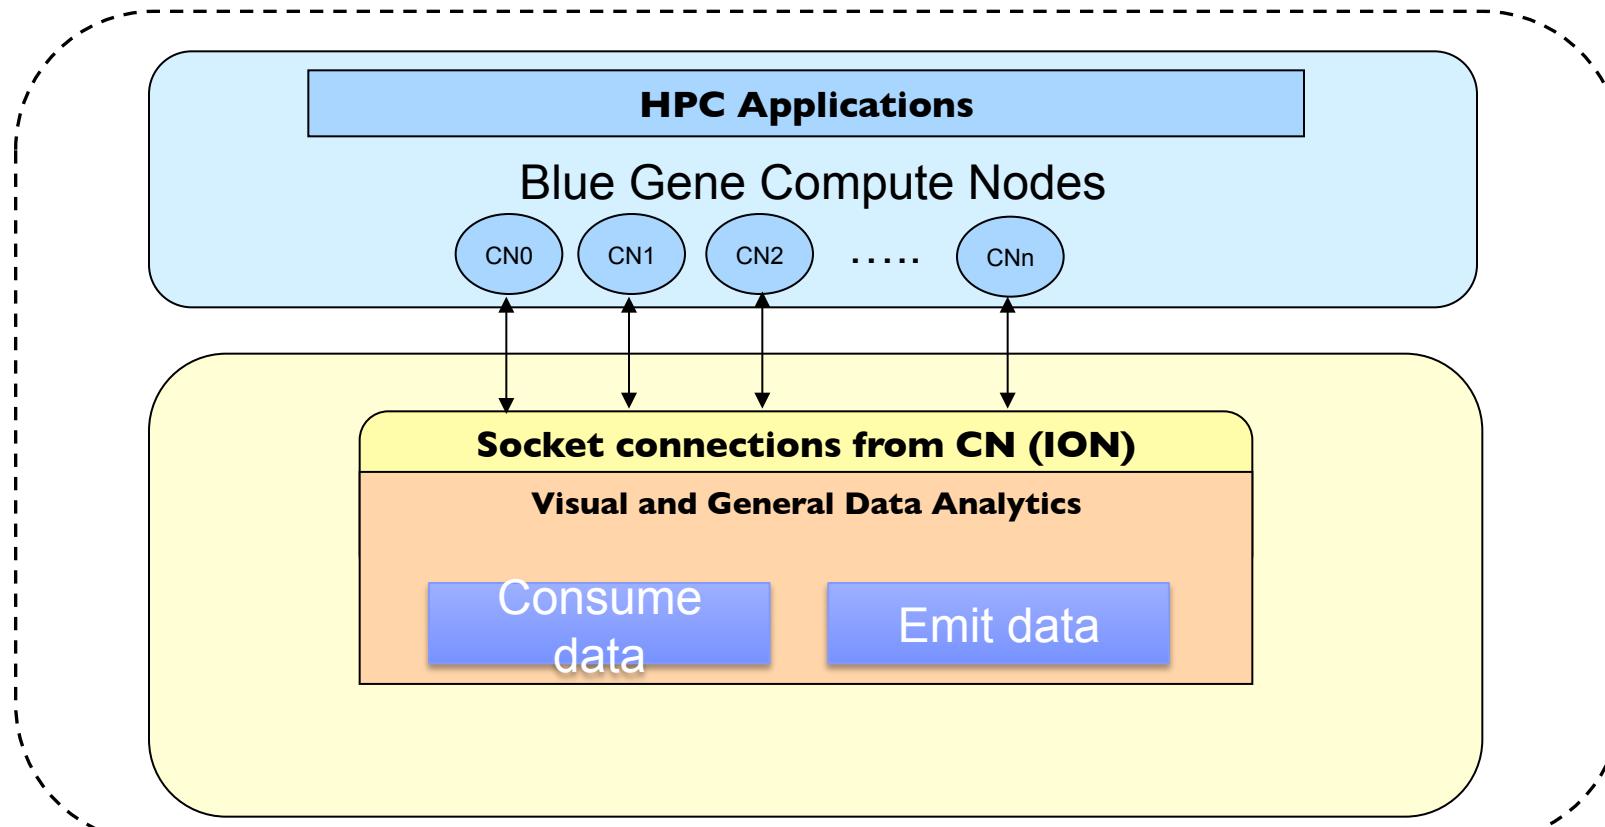
Connection Process

1. CN/ION on private network only accessible by front-end and service nodes
2. Relay server executes on FEN
3. Env variable passed to simulation with ipaddr:port of relay server
4. Simulation connects via socket
5. Client application connects to relay server via socket
6. Connection is established and XML encoded steerable parameters are transmitted from BG to client applications

Command/Parameter and Data Flow



Data Flow: HPC to Visualization Application



- Support different Data types
- Support 2-way flow
- Compute Node to Cluster Node (M:N, where M>=N)
- Visualization platform does NOT have to be co-located with HPC platform,
 - But there must be route from ION to visualization cluster, ie. Same subnet

Basic Application Enablement

Initialization

```
Steering_enable(REG_TRUE);  
Steering_initialize("TEST", numCommands, commands);
```

- **Register the output IO channels**

```
Register_IOType("TEST DATA",  
                REG_IO_OUT,  
                REG_IO_CLIENT,  
                1,  
                &(iotype_handle[0]))
```

- **Register a steerable parameter**

```
status = Register_param("MAX AMPLITUDE", REG_TRUE, (void *)(&amplitude), REG_INT, "1", "15");  
Finalize_registration();
```

Basic Application Enablement

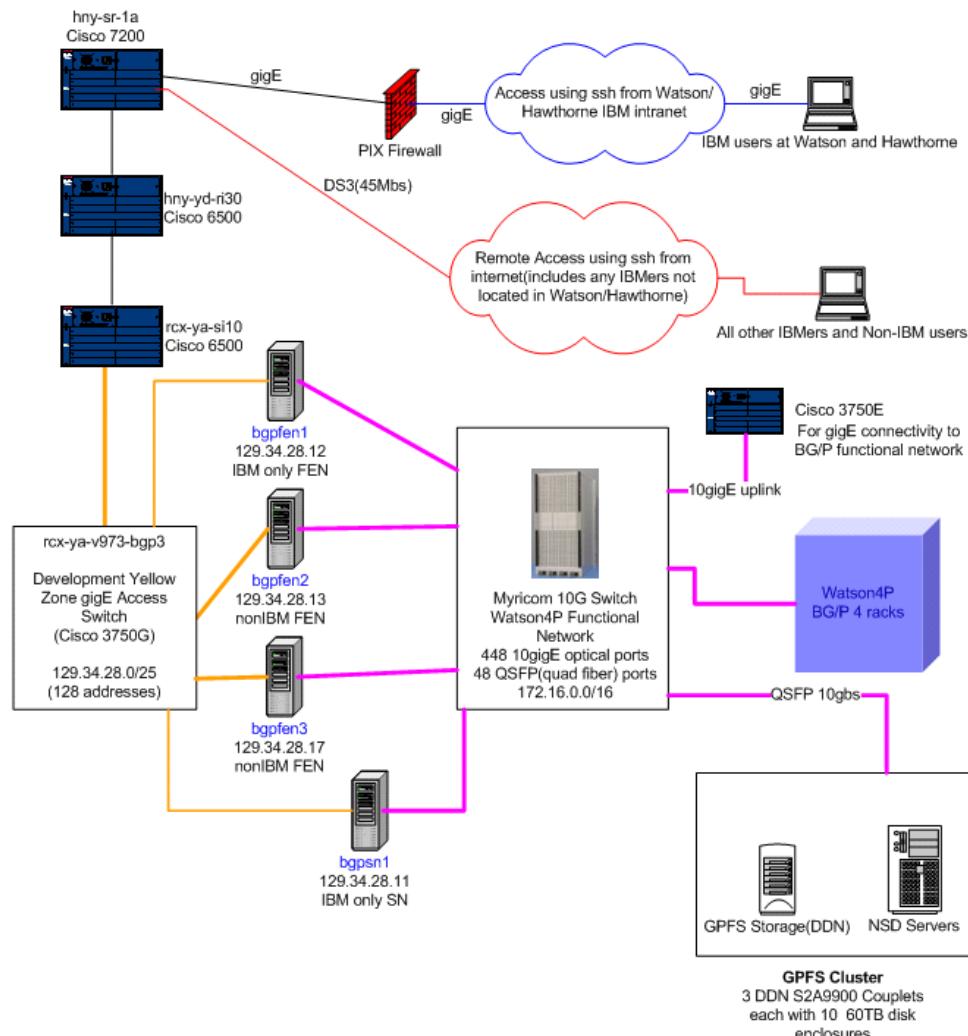
Main simulation loop

```
While simulating {  
    Steering_control(l, &num_params_changed, changed_param_labels,&num_recv_cmds,  
                    recv_cmds, recv_cmd_params);  
    Process recv_cmds  
    for all data to output {  
        Emit_start(iotype_handle[j], i, &iobhandle)  
        Emit_data_slice(iobhandle, REG_FLOAT, size, (void*) data_array);  
        Emit_stop();  
    }  
}
```

Note: To consume data you replace `Emit_start` with `Consume_start`, `Emit_data` with `Consume_data`, etc.

Watson4P Access and Functional Network Overview

Development Yellow Zone Access Design



Note: the Watson4P Service Network(10.0.0.0/16) is not shown

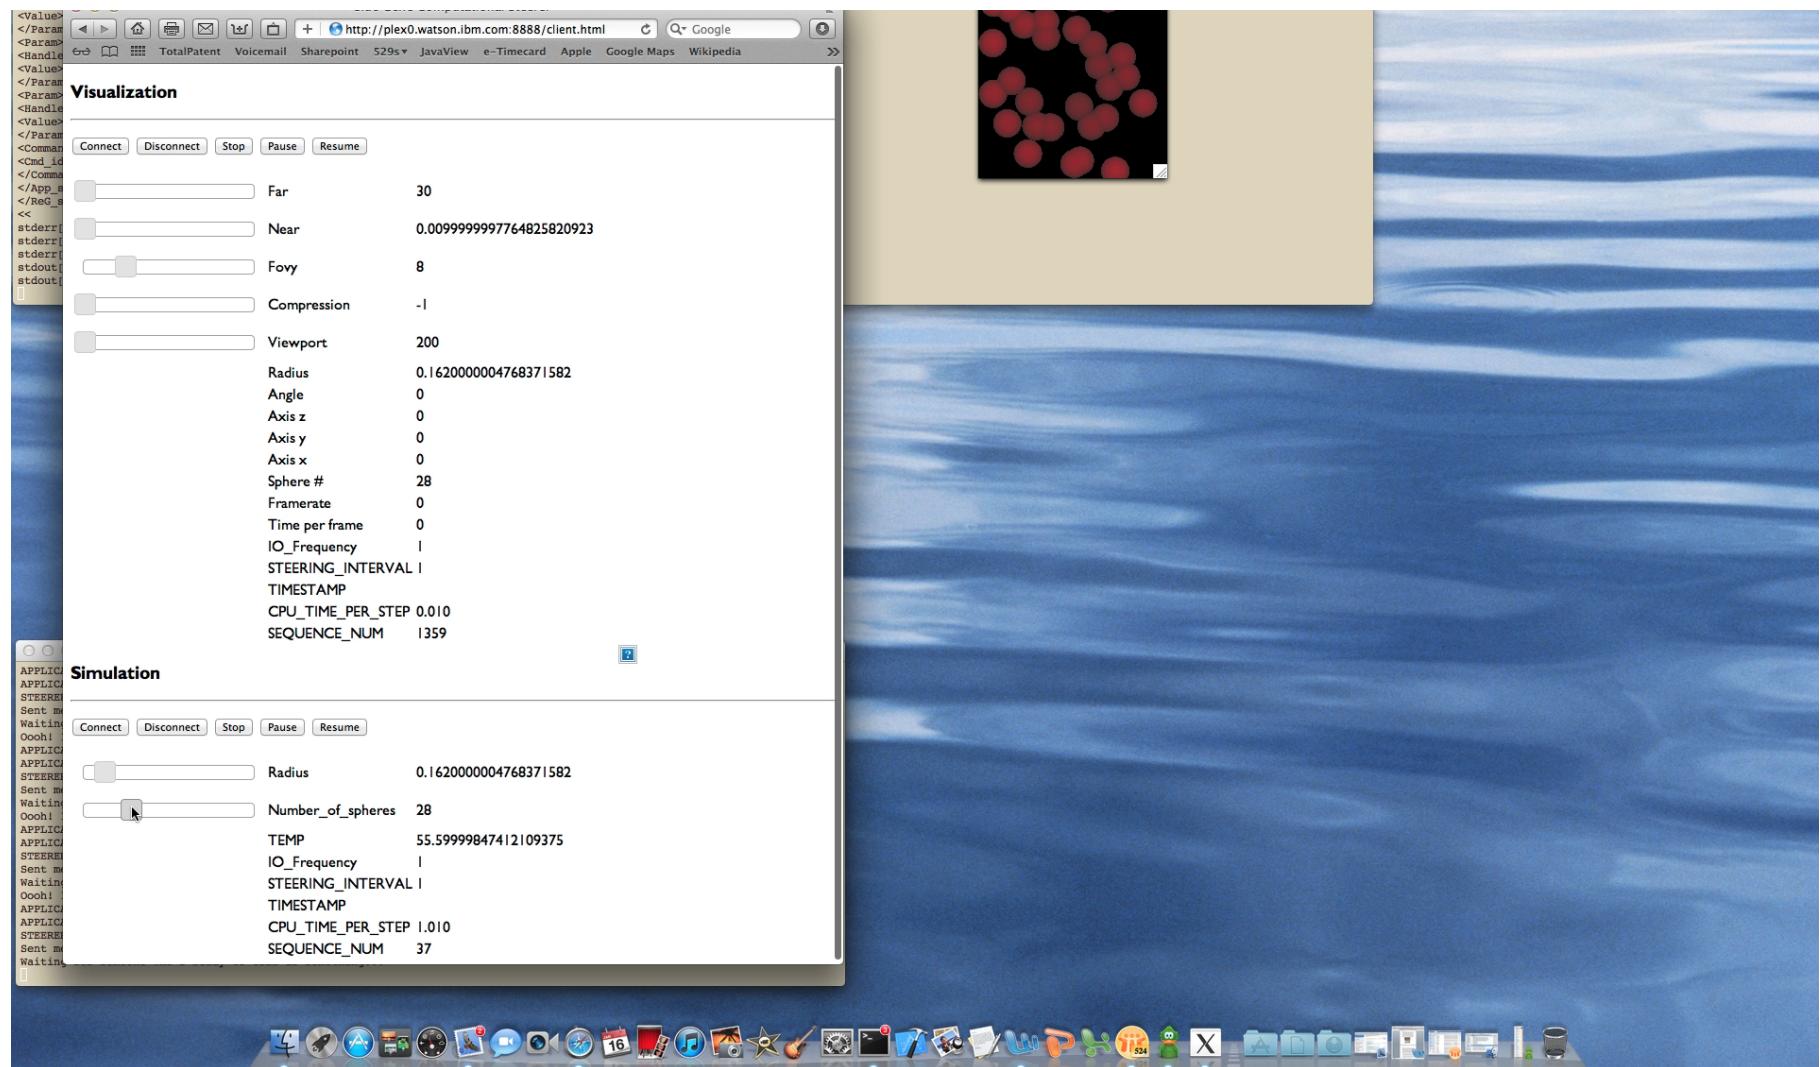
— Functional Network 10gigE
— Development Yellow Zone Network gigE

SN - Service Node
FEN - Front End Node

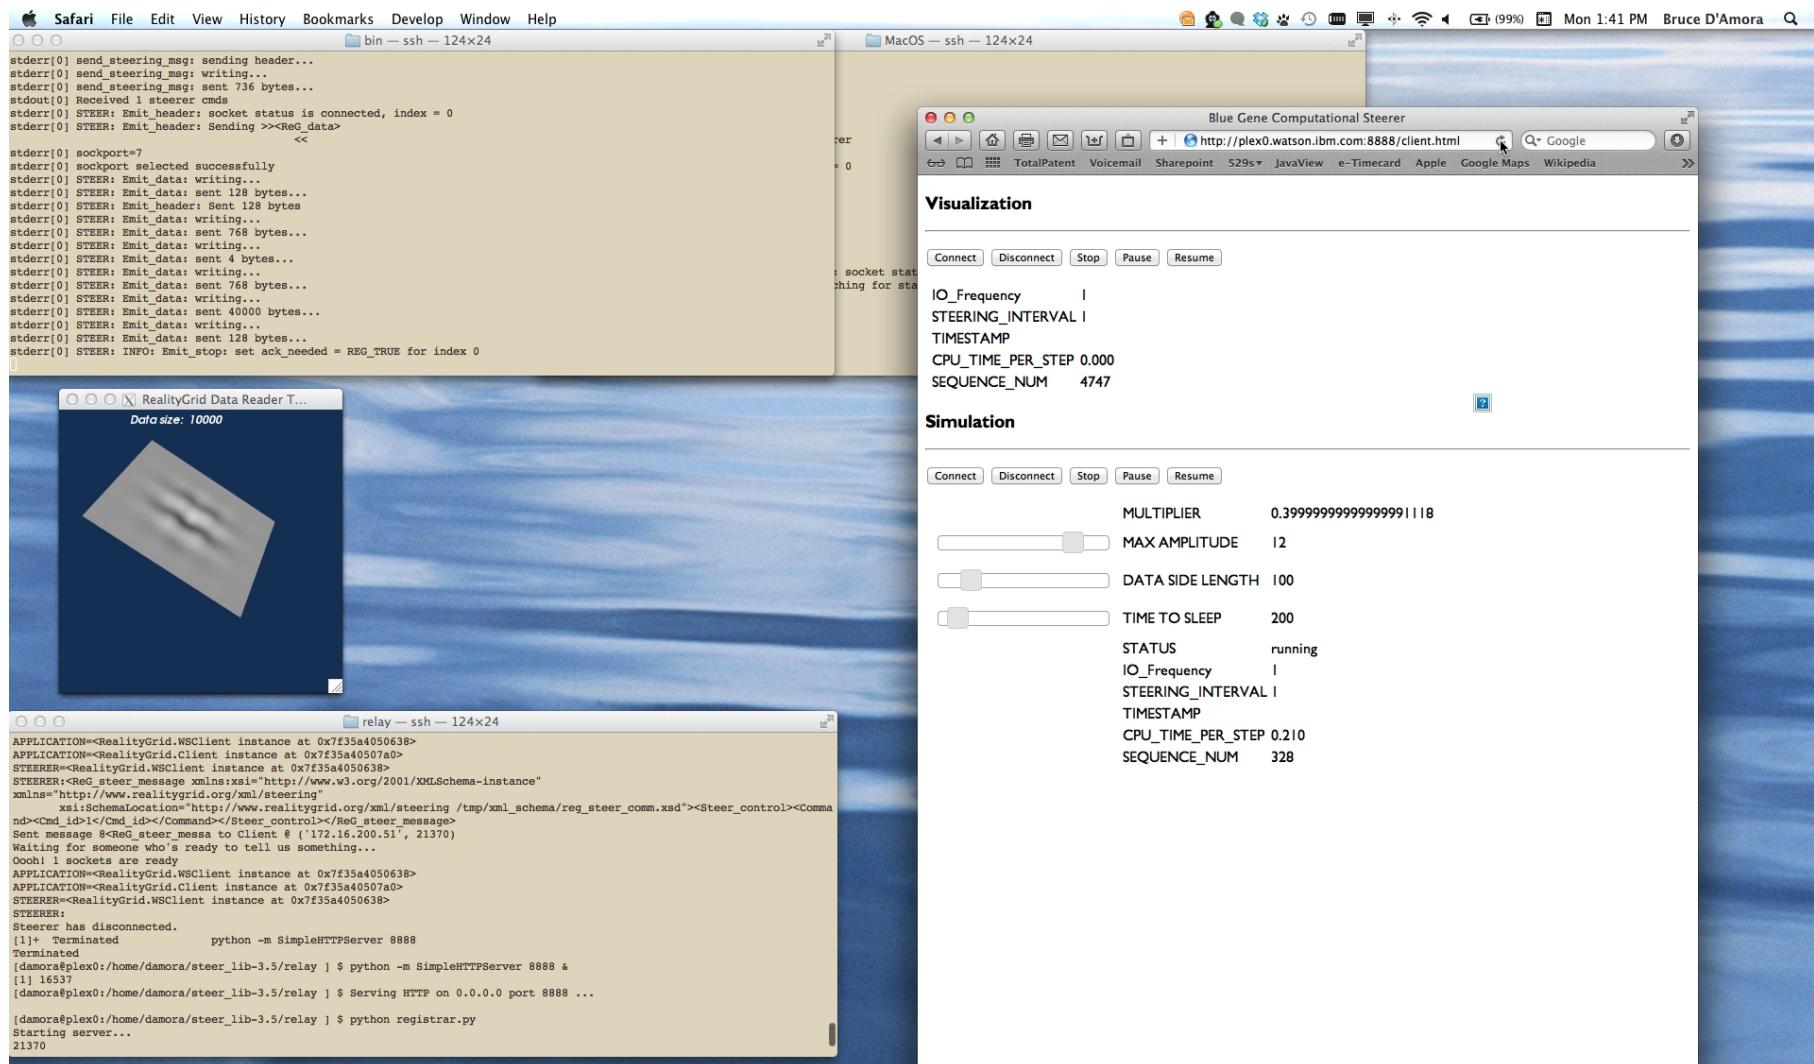
All nodes are
<name>.watson.ibm.com

Rev. 3 October 2011

Video



Video



New Steering Client

Computational Steering – Steering Client

User Info Available Steered Apps My Connected Steered Apps Admin Panel

A1328039724975854 A1328039725767912

Application Type :SIMULATION_APPS
Running on Hostname :127.0.0.1
Running on Port :1210

Monitored Parameters

Name	Value
CPU_TIME_PER_STEP	0.002
SEQUENCE_NUM	315
multiplier	0.6000000...

Update Interval(msec): 1000 Refresh

Steered Parameters

Name	Value	New Value
STEERING_INTERVAL	5	5
a string	running	
data side length	100	
max amplitude	12	

Tell Tell All Refresh

Data IO

Name	IO Type	Freq	New Freq
SOME_INPUT_DATA	INPUT	0	
mini_app visualization data	OUTPUT	1	

Consume Emit Tell Frequency

Checkpoint Types

Name	Freq	New Freq

Restart Create Checkpoint Tell Frequency

Message Log

Steering Users

Username
bdamora
bdanani

Send

This screenshot shows the 'Computational Steering – Steering Client' application interface. The main window title is 'Computational Steering – Steering Client'. The top menu bar includes tabs for 'User Info', 'Available Steered Apps', 'My Connected Steered Apps' (which is selected and highlighted in blue), and 'Admin Panel'. Below the tabs, two connection identifiers are listed: 'A1328039724975854' and 'A1328039725767912'. On the left, 'Application Type' is set to ':SIMULATION_APPS', running on 'Hostname :127.0.0.1' and 'Port :1210'. The 'Monitored Parameters' section lists CPU_TIME_PER_STEP (0.002), SEQUENCE_NUM (315), and multiplier (0.6000000...). The 'Steered Parameters' section shows STEERING_INTERVAL (5) and other parameters like a string (running), data side length (100), and max amplitude (12). The 'Data IO' section shows some_input_data as INPUT with freq 0 and mini_app visualization data as OUTPUT with freq 1. The 'Checkpoint Types' section is currently empty. The 'Message Log' section is a large empty text area. The 'Steering Users' section lists bdamora and bdanani. At the bottom right is a 'Send' button.

Next Steps

- **BlueGene/Q enablement**
 - Initial testing, but some bugs remain
- **Fully enable production-level application**
 - OpenFOAM is steerable, but no visualization of output yet
- **Performance characterization and optimization**

Acknowledgements

Thanks for the support of this work go to:

Dinesh Kaushik, David Keyes (KAUST)

Jim Sexton (IBM)